

Priority Scheduling

A presentation by Jakob H. and Laurin P.

Introduction

- every process has a priority
- task with the highest priority is ran first
- equal priority => first come first serve (FCFS)
- priority ranges not generally defined

An example

Assumption: lower number = higher priority

Process	Burst Time	Priority
P1	10	3
P2	1	1
P3	2	4
P4	1	5
P5	5	2

Resulting CPU time

Process	Burst Time	Priority
P1	10	3
P2	1	1
P5	5	2

Assigning Priorities

- External
 - manually setting a priority (nice)
 - how much did the customer pay?
 - how important is the task to the user?
- Internal
 - based on measurable quantities
 - number of open files
 - memory requirement
 - time limit

Preemptive vs Nonpreemptive

Preemptive

"preempt" CPU if priority of new process > running process

Nonpreemptive

Task will continue to run, new one is at head of queue

Starvation

- low-priority process never get to run
- higher priority processes "cut the line"
- two possible behaviours:
 1. process will run at an unconventional time
 2. process keeps waiting indefinitely

Starvation - a solution: aging

- increase priority of processes waiting for a long time
- for example:
 - priority from 127 (low) to 0 (high)
 - increase priority by 1 every 15 minutes
 - => process is ran after 32 hours at max
- avoids infinitely waiting processes (starvation)

When?

Advantages

- ease
- important tasks get more resources
- precise scheduling based on priority

Disadvantages

- crashes lead to processes being lost
- starvation (can be mitigated)

Thank you for your attention!